

Day 5

Teemu Sarapisto

May 22, 2018






1 Hands-on







- 1.1 Create the object skeleton with OpenCV's `ximgproc.thinning()` function from the original image and try if different values for the `thinningType` argument would make any difference

Main difference was that some of the thinners didn't create 1px wide lines.

1.2 Graphing







1.3 How long it took

around 4 hours

2 Homework 3

42% classification error with 10k samples. 52.8% classification error with 1k samples.

2.1 Code

```
#!/usr/bin/env python3
```

```

import pickle
import gzip
import matplotlib.cm as cm
import matplotlib.pyplot as plt
import numpy as np
import cv2

with gzip.open('mnist.pkl.gz', 'rb') as fs:
    train_set, valid_set, test_set = pickle.load(fs, encoding='latin1')

train_x, train_y = train_set
valid_x, valid_y = valid_set
test_x, test_y = test_set

def get_pixel_lbp(image, pixel, x, y):
    sum = 0
    if image[x-1, y-1] > pixel: sum += 1
    if image[x-1, y] > pixel: sum += 2
    if image[x-1, y+1] > pixel: sum += 4

    if image[x, y-1] > pixel: sum += 8
    if image[x, y+1] > pixel: sum += 16

    if image[x+1, y-1] > pixel: sum += 32
    if image[x+1, y] > pixel: sum += 64
    if image[x+1, y+1] > pixel: sum += 128
    return sum

def lbp_histogram(image):
    histogram = np.zeros(256)
    for y, row in enumerate(image):
        if y == 0 or y == image.shape[1] - 1: continue
        for x, cell in enumerate(row):
            if x == 0 or x == image.shape[0] - 1: continue
            histogram[get_pixel_lbp(image, cell, x, y)] += 1
    return histogram

def train(images):
    results = []
    for image in images:
        hist = lbp_histogram(image.reshape((28, 28)))
        results.append(hist)
    return results

count = 1000

```

```

results = train(train_x[:count])
labels = np.asarray(np.arange(0, count), dtype=np.float32).reshape(count, 1)
for idx, n in enumerate(results):
    labels[idx] = train_y[idx]
knn = cv2.ml.KNearest_create()

knn.train(np.asarray(results, dtype=np.float32), cv2.ml.ROW_SAMPLE, labels)

test_results = train(test_x[:count])
test_results = np.asarray(test_results, dtype=np.float32)

ret, predictions, neighbours, dist = knn.findNearest(test_results, k=1)

success = 0
for idx, res in enumerate(predictions):
    if int(res[0]) == test_y[idx]:
        success += 1

print(success / len(predictions))

```

2.2 How long it took

Around 4 hours. Though I was doing something really stupid with creating the labels for the KNN at first, had I not done that I could've easily done it in two hours...